

An Exact Hypothesis Test For Samples With Few Effective Clusters

AKIVA YONAH MEISELMAN*

This version: November 10, 2021

[Click here for most recent version](#)

I propose a hypothesis test for clustered samples. This test is exact in samples with few clusters, few ever-treated clusters, cluster size outliers, or treatment intensity outliers; these features cause previous tests to over- or under-reject true hypotheses. I derive my test by inverting the distribution of the test statistic under a standard assumption about the errors, so that critical values can be selected from a distribution that matches the test statistic. I use Monte Carlo simulations to demonstrate where this adjustment is most impactful in achieving exact tests compared to previous hypothesis tests, and I apply my test to an empirical setting from the literature.

Researchers often find that their samples include observations that are not independent. Rather, the observations are grouped into independent clusters. Common research designs based on fixed effects use the dependence between observations to avoid bias due to omitted unobservable characteristics. Statistical inference in such samples must account for clustering. Failing to do so can lead to dramatically inaccurate standard errors, confidence intervals, and p-values.

Since [White \(1984, theorem 6.3\)](#), there have been tools that allow researchers to perform cluster-robust inference. The tools are consistent; they are valid in

* University of Texas at Austin, Bernard and Audre Rapoport Building 3.134G, 2225 Speedway, Austin, TX 78705, Email: meiselman@utexas.edu.

asymptotically large samples. However, those tools are not guaranteed to work when the effective sample size is small. Conventional cluster-robust hypothesis tests will over-reject or under-reject true hypotheses in four environments: (1) when the number of clusters is small; (2) when the number of clusters with treatment variation is small; (3) when there are cluster size outliers; and (4) when there are treatment intensity outliers. Henceforth I refer to samples with one or more of these features as having few “effective clusters,” since the asymmetry between clusters in (2), (3), and (4) causes the test statistic to behave similarly to (1).

These environments with few effective clusters can easily come up in empirical work using common research designs. Consider the example of [Karaivanov et al. \(2021\)](#), a study of COVID-related mask mandates in Canada. Their panel included 13 Canadian provinces and territories, so the number of clusters was small. In [Myers \(2017\)](#), among 50 US states, only 4 states implemented an abortion-access policy she examined during the period of analysis. The number of clusters (states) with treatment variation (policy changed) was small. In a study of managed care and mortality in Texas, [Kuziemko, Meckel and Rossin-Slater \(2018\)](#) cluster by county, and over 40% of the population was located in 5 of the 254 counties. Those 5 counties are cluster size outliers. [Bound et al. \(2020\)](#) analyze the impact of state appropriations to public colleges on foreign public college enrollees. In that context, some states had much larger year-to-year changes in state appropriations, and those states are treatment intensity outliers. In all of these cases, some adjustment to the conventional cluster-robust hypothesis test would be necessary. Moreover, previous tests, including those recommended by the highly-cited [Cameron and Miller \(2015\)](#), are all vulnerable to one or more of these features.

In this paper, I propose an exact hypothesis test for clustered samples with cluster-level fixed effects that rejects true hypotheses at the correct rate even with few effective clusters. I develop my test by calculating the inverse of the

distribution of the conventional cluster-robust test statistic under the assumption that the errors are normal and homoskedastic. This is a common assumption made in the existing literature. Unlike previous hypothesis tests, my test uses the exact distribution of the test statistic under these assumptions. If the errors are not normal or homoskedastic, then my test performs similar to those in the literature; it is asymptotically consistent, in the sense that the rejection rate converges to the nominal size of the test. Thus, my test improves on the existing tests by performing as well as existing tests in large samples while also being an exact test in small samples under the required error conditions. The main advantage of my test over other asymptotically consistent tests is that my test performs better than those tests when there are few effective clusters.

I demonstrate using Monte Carlo simulations that my test is exact and that previous tests are not exact in samples with few effective clusters. I calculate the rejection rates of my test and the other tests in randomly-generated samples where I vary the number of clusters, the number of clusters with treatment variation, the size of a cluster size outlier, and the treatment intensity of a treatment intensity outlier. I also calculate rejection rates when the errors are non-normal, heteroskedastic, or serially correlated. My test performs as well as any other test when the assumption of normal, homoskedastic errors is violated. Additionally, I illustrate how my test functions in the setting from [Abouk and Adams \(2013\)](#), a study of the impact of texting-while-driving bans; in a simulation exercise based on their data, my test also performs well.

This is the first paper to explicitly target exact inference in a finite, clustered sample. There have been other studies which make adjustments to improve on the conventional cluster-robust hypothesis test, including [Hansen \(2007\)](#), [Bell and McCaffrey \(2002\)](#), [Carter, Schnepel and Steigerwald \(2017\)](#), and [Cameron, Gelbach and Miller \(2008\)](#), but those previous tests fail to reject at the correct rate when there are few effective clusters. Two of those studies, [Bell and McCaffrey \(2002\)](#) and [Carter, Schnepel and Steigerwald \(2017\)](#), make an assumption that is

similar to the one I use to derive my test (requiring normal, homoskedastic errors), but neither use the exact distribution of the test statistic, instead approximating the test-statistic with a t -distribution. My test does not use an approximation. As a result, my test performs better in samples with few effective clusters.

Section 1 gives context for my contribution to the literature on cluster-robust inference. Section 2 describes the model with clustering and cluster-level fixed effects that is the setting for this paper. Section 3 introduces my test and show that it is exact. Section 4 presents evidence from Monte Carlo simulations that my test rejects true hypotheses at the correct rate even when other tests fail to do so. Section 5 applies my test to the empirical setting in [Abouk and Adams \(2013\)](#), and Section 6 concludes.

1. Literature

This paper contributes to an existing literature that addresses cluster-robust inference with a particular focus on applied research designs. Consistent cluster-robust variance estimators (CRVEs) were developed by [White \(1984\)](#), [Liang and Zeger \(1986\)](#), and [Arellano \(1987\)](#). During the credibility revolution of the 1990s, research designs that included fixed effects came into much broader use. In the wake of this, [Bertrand, Duflo and Mullainathan \(2004\)](#) elevated awareness of CRVEs among applied economists.

I build on previous work that focuses on attaining asymptotically consistent inference in clustered samples. [White \(1984\)](#) shows that, when clusters are equally sized and the errors are homoskedastic, the basic CRVE (henceforth \hat{V}_{CR0}) can consistently estimate the variance of the OLS estimator. [Hansen \(2007\)](#) relaxes the homoskedasticity assumption, showing that equal-sized clusters alone allow the CRVE to be consistent and that it converges at a rate determined by the number of clusters G . He recommends critical values be drawn from a t -distribution with $G - 1$ degrees of freedom.

My proposed test is asymptotically consistent based on the same logic as

in [Carter, Schnepel and Steigerwald \(2017\)](#). They show that \hat{V}_{CR0} is consistent even when cluster sizes vary, but the rate of convergence is instead determined by G^* , what [Carter, Schnepel and Steigerwald \(2017\)](#) call the “effective number of clusters.” They recommend calculating G^* as a diagnostic tool. If G^* is large, inference can rely upon the asymptotic properties of the test statistic to determine its behavior, so critical values can reasonably be drawn from the standard normal distribution $N(0, 1)$. In this paper, I do not focus on actually calculating G^* , but I do use the phrase “few effective clusters” to refer to samples with small G^* .

I will compare my test to the hypothesis tests recommended by [Cameron and Miller \(2015\)](#), a paper that many applied economists use as a guide for how to handle cluster-robust inference. Those tests are derived from [Hansen \(2007\)](#), [Carter, Schnepel and Steigerwald \(2017\)](#), [Bell and McCaffrey \(2002\)](#), and [Cameron, Gelbach and Miller \(2008\)](#). Section 3.2 describes these tests in more detail; here, I give a brief overview.

[Hansen \(2007\)](#), [Carter, Schnepel and Steigerwald \(2017\)](#), and [Bell and McCaffrey \(2002\)](#) select critical values for the test statistic from a t -distribution. [Hansen \(2007\)](#) sets the degrees of freedom to $G - 1$ and [Carter, Schnepel and Steigerwald \(2017\)](#) sets the degrees of freedom to the effective number of clusters G^* . [Bell and McCaffrey \(2002\)](#) address finite-sample cluster-robust inference by developing two additional CRVEs (\hat{V}_{CR2} and \hat{V}_{CR3}), aimed at reducing bias in the variance estimation step. I build on their framework, making a similar but weaker assumption and discarding the approximation embedded in their test.

Taking a different approach, [Cameron, Gelbach and Miller \(2008\)](#) generate a reference distribution for the test statistic through a resampling method, the wild cluster bootstrap with restricted residuals (henceforth WCR)¹. [Djogbenou, MacKinnon and Nielsen \(2019\)](#) show that, in addition to performing well in simulations, WCR is a formal asymptotic refinement of the conventional test based

¹The “restricted residuals” used in WCR are the residuals from the model estimated subject to the restriction that the null hypothesis is true.

on \hat{V}_{CR0} and Hansen's $G - 1$ degrees of freedom.

2. Model

In this section, I describe a linear model with clustering in a single dimension and cluster-level fixed effects. I include fixed effects because they are a common feature of models where there may also be concerns about clustering. Consider the model:

$$(1) \quad y_{ig} = x_{ig}\beta + \gamma_g + \epsilon_{ig}$$

where x_{ig} is a $(1 \times K)$ vector of K covariates and γ_g is a cluster-level fixed effect. Clusters are indexed by g , and individual observations are indexed by i . Let N_g be the (deterministic) number of observations in cluster g , and then $N = \sum_g N_g$ is the total number of observations. Additionally, for ease of notation:

$$Y_g = \begin{bmatrix} y_{1,g} \\ y_{2,g} \\ \dots \\ y_{N_g,g} \end{bmatrix}, \quad X_g = \begin{bmatrix} x_{1,g} \\ x_{2,g} \\ \dots \\ x_{N_g,g} \end{bmatrix}, \quad \epsilon_g = \begin{bmatrix} \epsilon_{1,g} \\ \epsilon_{2,g} \\ \dots \\ \epsilon_{N_g,g} \end{bmatrix}$$

And similarly let Y (an $(N \times 1)$ matrix), X (an $(N \times k)$ matrix), and ϵ (an $(N \times 1)$ matrix) stack up the outcomes, covariates, and errors of all the clusters, so that X_g contains the rows of X corresponding to cluster g .

For standard fixed-effects estimation of β , the fixed effects γ_g are absorbed. That is, rather than estimate γ_g , I transform the sample using cluster-level averages². Let $\ddot{Y}_g = Y_g - \frac{1}{N_g} \sum_{i=1}^{N_g} y_{ig}$, $\ddot{X}_g = X_g - \frac{1}{N_g} \sum_{i=1}^{N_g} x_{ig}$, and $\ddot{\epsilon}_g = \epsilon_g - \frac{1}{N_g} \sum_{i=1}^{N_g} \epsilon_{ig}$. Assuming that $\mathbb{E}(\epsilon_g \mid X_g) = 0$, the fixed effects estimator

²For an extensive treatment of fixed effects models and fixed effects absorption (also called the “within transformation”), including the consistency and unbiasedness of the fixed effects estimator, see [Hansen \(2021\)](#), sections 17.8, 17.9, and 17.20.

can consistently estimate β :

$$(2) \quad \hat{\beta} = (\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T \ddot{Y}$$

For inference on β , I examine two-sided tests of hypotheses with the form $H_0 : c_0^T \beta = a_0$, where c_0 is a $(K \times 1)$ vector and a_0 is a scalar. Without loss of generality, I normalize³ c_0 so that $c_0^T c_0 = 1$. Inference involves calculating a test statistic t_0 and comparing it to a critical value q^* . An exact test will reject a true hypothesis with some probability α , which is the “size” of the test.

Calculating t_0 begins with estimating $\hat{V}(\hat{\beta})$. The true variance of $\hat{\beta}$ is given by:

$$V(\hat{\beta}) = (\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T \mathbb{E}(\ddot{\epsilon} \ddot{\epsilon}^T) \ddot{X} (\ddot{X}^T \ddot{X})^{-1}$$

In a sample of independent, identically distributed observations, inference could rely on the assumption that the errors are all mutually independent. However, in this clustered setting, I make only the (standard) weaker assumption that the errors are uncorrelated across clusters:

$$\mathbb{E}(\epsilon_g \epsilon_{g'}^T) = 0, \forall g \neq g'$$

Let $\hat{\epsilon}_g = \ddot{Y}_g - \ddot{X}_g \hat{\beta}$ be the residuals for cluster g . The simplest cluster-robust variance estimator, \hat{V}_{CR0} from [White \(1984\)](#), takes the form:

$$\hat{V}(\hat{\beta}) = (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T \hat{\epsilon}_g \hat{\epsilon}_g^T \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1}$$

Finally, a test statistic can be generated, using the parameter estimator $\hat{\beta}$ and

³Suppose there is some hypothesis such that $\tilde{c}_0^T \tilde{c}_0 \neq 1$. Note that $\tilde{c}_0^T \tilde{c}_0 > 0$, since $\tilde{c}_0^T \tilde{c}_0 = 0$ would not actually be testing a linear hypothesis. Then let $c_0 = \frac{\tilde{c}_0}{\sqrt{\tilde{c}_0^T \tilde{c}_0}}$, so that $c_0^T c_0 = \frac{\tilde{c}_0^T \tilde{c}_0}{\tilde{c}_0^T \tilde{c}_0} = 1$.

the variance estimator $\hat{V}(\hat{\beta})$:

$$t_0 = \frac{c_0^T \hat{\beta} - a_0}{\sqrt{c_0^T \hat{V}(\hat{\beta}) c_0}}$$

If the hypothesis H_0 is true, then generating a test statistic with a large magnitude is relatively unlikely. So t_0 can be compared to some critical value q^* , and H_0 is rejected if $|t_0| > q^*$.

3. Hypothesis Tests

3.1. My Proposed Test

I propose a new method of testing linear hypotheses, which I develop here. I prove that my test is valid in two cases: normal, homoskedastic errors; and asymptotically large samples.

When the errors are normal and homoskedastic (specifically, see Assumption 1 below), my test rejects true hypotheses with probability equal to the nominal test size α . This assumption is standard among finite-sample adjustments to cluster-robust inference. Furthermore, under the (weaker) assumptions that allow all cluster-robust hypothesis tests to be consistent, my test is also consistent in the sense that its rejection rate converges in probability to α . In other words, my test maintains the good asymptotic properties of previous cluster-robust hypothesis tests.

In order to perform an exact hypothesis test, I would like to find a critical value $q^*(H_0, \alpha)$ such that, if H_0 is true, then:

$$P(|t_0| > q^*(H_0, \alpha)) = \alpha$$

The optimal method for selecting critical values would be some $q^*(\cdot, \cdot)$ that gives an exact test for any hypothesis H_0 and any test size α .

If the CDF of t_0^2 were known, I could invert that CDF to find the critical

value. Using $F_{t_0^2}(\cdot)$ to denote the CDF of t_0^2 , that critical value $q^*(\cdot, \cdot)$ is given by:

$$\begin{aligned} F_{t_0^2}((q^*(H_0, \alpha))^2) &= 1 - \alpha \\ q^*(H_0, \alpha) &= \sqrt{F_{t_0^2}^{-1}(1 - \alpha)} \end{aligned}$$

The intuition for my test is that I make an assumption that is strong enough to determine $F_{t_0^2}(\cdot)$.

ASSUMPTION 1: *The errors are normal and homoskedastic with an intraclass correlation coefficient ρ that is constant within and across clusters:*

$$\epsilon_g \sim N(0, \sigma^2 \Omega_g), \quad \Omega_g = \begin{bmatrix} 1 & \rho & \dots & \rho \\ \rho & 1 & \dots & \rho \\ \dots & \dots & \dots & \dots \\ \rho & \rho & \dots & 1 \end{bmatrix}$$

Note that σ and ρ may be unknown. I refer to this assumption as “standard” because it is a slightly weaker version of the same assumption that several other papers use to adjust cluster-robust inference in finite samples ([Carter, Schnepel and Steigerwald, 2017](#); [Bell and McCaffrey, 2002](#)). With this assumption, $F_{t_0^2}(\cdot)$ can be derived in Theorem 1 below.

I introduce some notation in preparation for the test definition. Let I_g be an identity matrix of size N_g , and let ι_g be a column vector of length N_g whose elements are all 1. So then let $M_g = I_g - \frac{1}{N_g} \iota_g \iota_g^T$. Additionally, let $H = \ddot{X}(\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T$, let I be an identity matrix of size N , and let $(I - H)_g$ be the rows of $(I - H)$ corresponding to cluster g . Furthermore, I define a function $L(\cdot; X, H_0)$ that I will show is equal to the CDF $F_{t_0^2}(\cdot)$ in Theorem 1:

$$(3) \quad \text{Let } L(q; X, H_0) = \frac{1}{2} - \frac{1}{\pi} \int_0^\infty \frac{\sin\left(\frac{1}{2} \sum_{j=1}^{G+1} \tan^{-1}(\lambda_j u)\right)}{u \prod_{j=1}^{G+1} \left(1 + \lambda_j^2 u^2\right)^{\frac{1}{4}}} du$$

where $d_0 = \ddot{X}(\ddot{X}^T \ddot{X})^{-1} c_0$, $d_g = (I - H)_g^T A_g^T \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0$,
 $D_+ = [d_0 \quad d_1 \dots d_g \dots d_G]$, $D_- = [\frac{1}{q} d_0 \quad -d_1 \dots -d_g \dots -d_G]$,
 M block-diagonal, with g -th block M_g , and
 (λ_j) are the eigenvalues of $D_-^T M D_+$

To implement my test⁴:

- 1) Calculate the test statistic $t_0 = \frac{c_0^T \hat{\beta} - a_0}{\sqrt{c_0^T \hat{V}(\hat{\beta}) c_0}}$
- 2) Find $L(q; X, H_0)$
- 3) Determine the critical value q^* such that $L((q^*)^2; X, H_0) = 1 - \alpha$
- 4) Reject H_0 if $|t_0| > q^*$

Note that, since $L(q; X, H_0)$ is increasing in q , $L(\cdot; X, H_0)$ can easily be inverted numerically. Finally, let $\alpha^*(X, H_0) = P(|t_0| > q^*(\alpha; X, H_0))$ be the rejection rate of my test – the rate at which my test rejects a true hypothesis H_0 .

THEOREM 1: *If Assumption 1 holds and H_0 is true, then:*

$$F_{t_0^2}(q) = L(q; X, H_0)$$

and $\alpha^* = \alpha$

where $L(\cdot; \cdot, \cdot)$ is defined in Equation 3.

I prove Theorem 1 by deriving $L(\cdot; \cdot, \cdot)$ in Appendix A. In the simplest version of the test, I calculate the test statistic t_0 using the variance estimator \hat{V}_{CR0} . In Section 4, I use two additional variants of my test that are based on different variance estimators. These variants use \hat{V}_{CR2} and \hat{V}_{CR3} , the estimators given by Bell and McCaffrey (2002). The proof of Theorem 1 in Appendix A holds for both of these variants; they are exact tests under Assumption 1.

⁴My test is available in R as the function “p.value.meis()” in the package “clubsoda”, available through github.

My test using \hat{V}_{CR0} is also asymptotically consistent under the relatively weak assumptions described by [Carter, Schnepel and Steigerwald \(2017\)](#). They demonstrate that, when using \hat{V}_{CR0} , the test statistic converges to a standard normal distribution: $t_0 \xrightarrow{d} N(0, 1)$. I build on this result, using their two main assumptions. The first ensures that the errors have finite fourth moments. The second ensures that the observations aren't too concentrated in a small number of clusters.

ASSUMPTION 2: *For each cluster g , there is some positive scalar B and some $(N_g \times N_g)$ matrix Ω_g such that $\epsilon_g = \Omega_g^{\frac{1}{2}} \eta_g$, where the η_g is a vector of length N_g whose elements are uncorrelated random variables and where:*

$$\begin{aligned} \mathbb{E}(\eta_{ig}\eta_{jg}\eta_{kg}\eta_{lg}) &= 0, & \mathbb{E}(\eta_{ig}\eta_{jg}\eta_{kg}^2) &= 0 \\ \mathbb{E}(\eta_{ig}\eta_{jg}^3) &= 0, & \mathbb{E}(\eta_{ig}^2\eta_{jg}^2) &= 0 \\ \mathbb{E}(\eta_{ig}^4) &\leq B \end{aligned}$$

ASSUMPTION 3: *Let $\lambda_g^{CSS} = c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T \mathbb{E}(\ddot{\epsilon}_g \ddot{\epsilon}_g^T) \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0$, let $V_0 = c_0^T V(\hat{\beta}) c_0$, and let $P_g = (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T \ddot{X}_g$. As $n \rightarrow \infty$:*

$$\begin{aligned} G &\rightarrow \infty \\ \mathbb{E} \left(\frac{\sum_g (\lambda_g^{CSS})^2}{\left(\sum_g \lambda_g^{CSS}\right)^2} \right) &\rightarrow 0 \\ \frac{1}{V_0} c_0^T \left(\sum_g \left(P_g - \frac{1}{G} I \right) V(\hat{\beta}) \left(P_g - \frac{1}{G} I \right)^T \right) c_0 &\xrightarrow{p} 0 \end{aligned}$$

Recall that α^* is the rejection rate of my test.

THEOREM 2: *If Assumptions 2 and 3 hold and H_0 is true, then:*

$$\alpha^* \xrightarrow{p} \alpha$$

I prove Theorem 2 in Appendix B. The intuition for the proof is that the test statistic t_0 and the reference distribution defined by my test both converge in distribution to $N(0, 1)$. In large samples, my test will reject a true hypothesis with probability α . Proving that the variants of my test which use \hat{V}_{CR2} and \hat{V}_{CR3} are also asymptotically consistent remains an area of future work.

I have shown that my test's rejection rate α^* is equal to the nominal test size α in finite samples under Assumption 1 (normal, homoskedastic errors). I have shown separately, without Assumption 1, that α^* converges to α in an asymptotically large sample under Assumptions 2 and 3, so that my test is approximately exact when there are many effective clusters. Section 4 uses Monte Carlo simulations to demonstrate that this test performs well (rejecting at roughly the correct rate) even when dealing with few effective clusters and even when Assumption 1 does not hold.

3.2. Other Tests

In this section, I briefly discuss how previous cluster-robust hypothesis tests from the literature work and how they relate to my tests. Specifically, I will look at the tests recommended in Cameron and Miller (2015). These tests can be roughly divided into analytic tests, which select a critical value for t_0 from a known distribution, and resampling-based tests, which generate a simulated distribution of test statistics from which critical values are drawn. In Section 4, I compare my test's performance with the performance of these other tests.

In a test I refer to as “Hansen”, derived from Hansen (2007), it is recommended to estimate $\hat{V}(\hat{\beta})$ with \hat{V}_{CR3} and to select critical values for the test statistic t_0 from $T(G - 1)$, a t -distribution with $G - 1$ degrees of freedom.

The test from Bell and McCaffrey (2002), henceforth “BM”, involves estimating $\hat{V}(\hat{\beta})$ with \hat{V}_{CR2} and selecting critical values for t_0 from $T(m)$, where m is calculated according to a “Satterthwaite approximation” of t_0 . For an explanation of how the Satterthwaite approximation works and what assumptions it

relies on, see Appendix C.

Cameron and Miller (2015) also recommend a test, henceforth “CSS”, derived from Carter, Schnepel and Steigerwald (2017). In this test, $\hat{V}(\hat{\beta})$ is estimated with \hat{V}_{CR0} , and critical values for t_0 are selected from $T(G^*)$, where G^* is called the “effective number of clusters”. In Appendix C, I show how G^* is a simplified version of the Satterthwaite approximation.

Hansen, BM, and CSS all approximate the test statistic t_0 as a t -distribution. By contrast, the last method recommended by Cameron and Miller (2015) is a resampling method, the wild cluster bootstrap. Using this method, $\hat{V}(\hat{\beta})$ is estimated with \hat{V}_{CR0} , and then over many bootstrap iterations, the residuals are resampled by multiplying them by values drawn from an auxiliary distribution with mean 0 and variance 1. A critical value t^* is then selected from the bootstrapped distribution of test statistics.

For choosing among the various specifications of the wild cluster bootstrap, I follow Djogbenou, MacKinnon and Nielsen (2019), a more recent study that tested many variants in simulations. They recommend resampling the restricted residuals (the residuals from the restricted model, subject to H_0), with the auxiliary distribution being either the Rademacher distribution or the Mammen distribution. I refer to these tests as “WCR-R” and “WCR-M”, respectively.

There are some parallels between my test and previous analytic tests in the literature. However, these other methods all approximate the distribution of the test statistic. By contrast, I have made an assumption that is strong enough to fully determine the distribution of the test statistic. In the next section, I will demonstrate that this approach makes my test perform better in many samples; my test rejects true hypotheses at the correct rate even when other tests fail to do so.

4. Simulations

In this section, I show the results of Monte Carlo simulations that demonstrate that my test is exact and that previous tests fail to reject true hypotheses at the correct rate under various conditions. Specifically, I focus on samples with few effective clusters: few clusters, few clusters with treatment variation, cluster size outliers, and treatment intensity outliers. I borrow certain features of the data-generating process from [Djogbenou, MacKinnon and Nielsen \(2019\)](#), another simulation study of cluster-robust inference.

Depending on the specification, I vary certain features of the design matrix X :

- G = number of clusters
- J = number of clusters with treatment variation
- N_1 = size of first cluster
- ϕ = treatment intensity of first cluster

In my simulation experiments, I use the following data-generating process:

$$\begin{aligned}
 (4) \quad x_{1ig} &= \mathbb{1}(g \leq J) \times \phi^{\mathbb{1}(g=1)} \times \frac{x_{1ig}^* - 8}{4}, \quad x_{1ig}^* \sim \chi_8^2 \\
 x_{2ig} &= \frac{x_{2ig}^* - 8}{4}, \quad x_{2ig}^* \sim \chi_8^2 \\
 y_{ig} &= \beta_0 + \beta_1 x_{1ig} + \beta_2 x_{2ig} + \gamma_g + \epsilon_{ig}
 \end{aligned}$$

Following [Djogbenou, MacKinnon and Nielsen \(2019\)](#), the covariates x_{1ig} and x_{2ig} are generated with distributions that are both skewed and leptokurtic. This highlights the fact that my main result does not require normally-distributed covariates.

For each cluster beyond the first, there are 5 observations in that cluster — that is, for $g > 1$, $N_g = 5$. I set $\beta_0 = 1$, $\beta_1 = 2$, and $\beta_2 = 3$. Since fixed effects are absorbed before any estimation, the values of γ_g do not affect estimation or


inference, so for simplicity I set $\gamma_g = 0$ for all g . Unless otherwise noted, the specification parameters have default values $G = 200$, $J = 200$, $N_1 = 5$, and $\phi = 1$.

In this section, I generate $\epsilon_{ig} \sim N(0, 1)$; this data generating process meets the conditions of Assumption 1. In Section 4.1, I will alter this specification with several violations of Assumption 1. Besides simply confirming what I showed in Theorem 1, this set of simulations serves to demonstrate the conditions where previous tests tend to over- or under-reject true hypotheses.

In each simulation, I generate a sample according to the data generating process, and I estimate $\hat{\beta}$ with the standard fixed effects estimator. Then, I test the true hypothesis $H_0 : \beta_1 = 2$ using my test as well as each of the tests recommended by Cameron and Miller (2015). Different tests require estimating $\hat{V}(\hat{\beta})$ using different variance estimators and comparing test statistics to critical values selected according to different methods. I do this for each of the different tests discussed in Section 3.

First, I present results for simulations in which the number of clusters G takes on the following values: $G = 5, 10, 20, 50, 100, 200$. In Figure 1, I plot the rejection rates in this DGP for each of the tests discussed in Section 3. In general, the degree of a given test's over- or under-rejection depends on the size of the test α . I therefore show rejection rates for 5% tests in Panel 1a and 1% tests in Panel 1b. As the number of clusters gets small, CSS tends to overreject and WCR-M tends to underreject fairly dramatically. Hansen and BM look more reasonable but also tend to underreject for small G . WCR-R rejects at the correct rate except when $G = 5$, where it seems to fail completely. When a sample has a small number of clusters, my test is the only exact test.

Next, in Figure 2, I show rejection rates from simulations where I vary J , the number of clusters with treatment variation. As seen in (4), for clusters beyond the J -th cluster, I simply multiply the value of x_{1ig} by 0. Hansen and CSS both overreject and WCR-M underrejects at $J \leq 20$. BM and WCR-R both struggle at

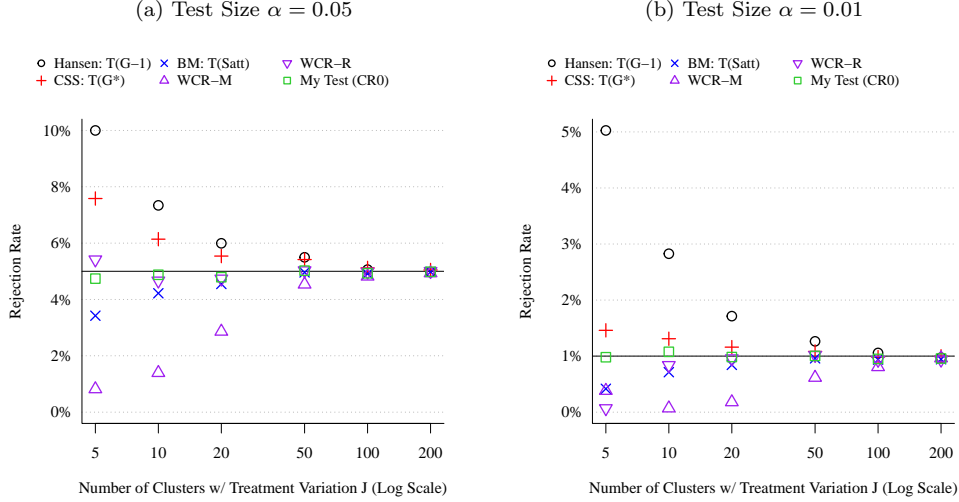
Figure 1. : Comparison of Hypothesis Tests With Varying G 

Notes: In these simulations, $J = 200$, $N_1 = 5$, and $\phi = 1$. The errors are i.i.d. standard normal. Each specification has 30,000 simulations.

$J = 5$. When a sample has a small number of clusters with treatment variation, my test is the only exact test.

Figure 3 plots rejection rates for different degrees of cluster size heterogeneity, where $N_1 = 20, 100, 200, 500$. In each specification, there are 200 clusters, and for $g > 1$, $N_g = 5$. So when $N_1 = 500$, the first cluster contains about a third of the observations in the sample. In that most extreme case, BM and CSS underreject, Hansen overrejects, and WCR-M overrejects for a 5% test only. My test is exact here, and WCR-R seems to at least be resilient to this form of cluster size heterogeneity.

In Figure 4, I show results for several values of the treatment intensity outlier parameter, so $\phi = 1, 5, 9, 13, 18, 24, 30$. Recall that the value of x_{1ig} is multiplied by ϕ for $g = 1$ only, so that a large value of ϕ creates a cluster that is an outlier in terms of variance in the treatment variable. Hansen overrejects for $\phi \geq 9$ and BM and CSS both underreject for $\phi \geq 9$. Around $\phi = 18$, both WCR-R and


Figure 2. : Comparison of Hypothesis Tests With Varying J 

Notes: In these simulations, $G = 200$, $N_1 = 5$, and $\phi = 1$. The covariate x_{1ig} is distributed as $x_{1ig} = \mathbb{1}(g \leq J) \times \phi^{\mathbb{1}(g=1)} \times \frac{x_{1ig}^* - 8}{4}$, where $x_{1ig}^* \sim \chi_8^2$. The errors are i.i.d. standard normal. Each specification has 30,000 simulations.

WCR-M begin to substantially underreject. In the presence of a large treatment intensity outlier, my test is the only exact test.

As discussed in Section 3, several of the methods (Hansen, CSS, and BM) approximate the test statistic t_0 with a t -distribution. Of those, Hansen and CSS can underreject or overreject, depending on the specification. This happens because variance estimation and critical value selection lead to different biases. The approximation to a t -distribution depends implicitly on having an unbiased variance estimator⁵. However, Hansen uses \hat{V}_{CR3} , which is biased up in this DGP (corresponding to underrejection), and CSS uses \hat{V}_{CR0} , which is biased down in this DGP (corresponding to overrejection). It is also true that both of these methods can select inappropriate critical values due to the approximation itself. The reason that Hansen and CSS underreject in some specifications and overreject

⁵Since my test does not approximate the test statistic as a t -distribution, it does not require unbiased variance estimation

Figure 3. : Comparison of Hypothesis Tests With Varying N_1 

Notes: In these simulations, $G = 200$, $J = 200$, and $\phi = 1$. The errors are i.i.d. standard normal. Each specification has 30,000 simulations.

in other specifications is that the bias in the variance estimator causes the rejection rate to move in one direction and misspecified critical value selection causes the rejection rate to move in the opposite direction. For a deeper discussion of the approximation of the test statistic t_0 with a t -distribution, see Section C.

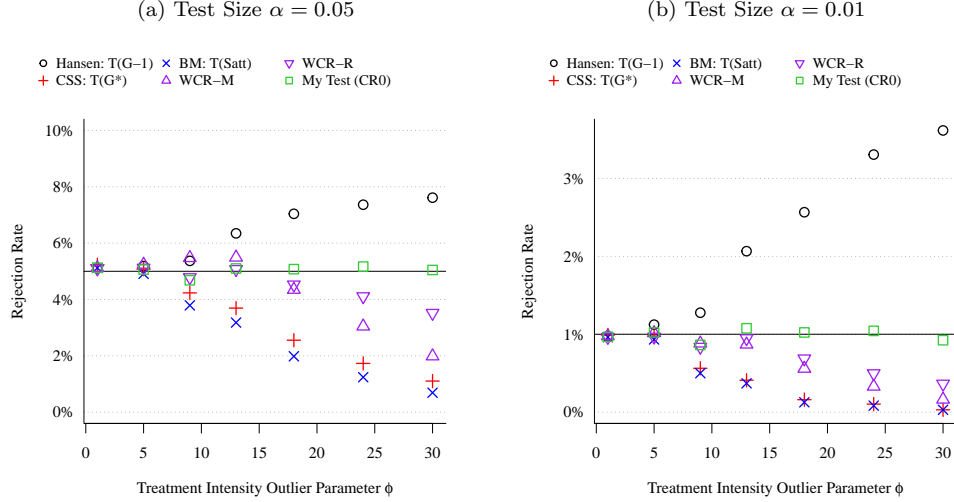
4.1. Robustness

So far, the Monte Carlo simulations have met the conditions of Assumption 1. Here, I present additional simulation evidence regarding the robustness of my test to violations of Assumption 1.

Since my test is asymptotically consistent, violations of Assumption 1 can only affect the performance of my test in samples with few effective clusters.

I focus on three violations of Assumption 1:

- Non-normal errors
- Serial correlation, where ϵ_{ig} is an AR(1) process

Figure 4. : Comparison of Hypothesis Tests With Varying ϕ 

Notes: In these simulations, $G = 200$, $J = 200$, and $N_1 = 5$. The covariate x_{1ig} is distributed as $x_{1ig} = \mathbb{1}(g \leq J) \times \phi^{\mathbb{1}(g=1)} \times \frac{x_{1ig}^* - 8}{4}$, where $x_{1ig}^* \sim \chi_8^2$. The errors are i.i.d. standard normal. Each specification has 30,000 simulations.

- Heteroskedasticity

These three violations correspond roughly to different parts of Assumption 1: normality, constant intraclass correlation, and homoskedasticity. It may be that normality of the errors can be relaxed when N_g , the number of observations per cluster, is large, and proving this is an area for future work. Still, I test robustness to non-normal errors here. [Bertrand, Duflo and Mullainathan \(2004\)](#) highlight serial correlation as an important potential problem in differences-in-differences analyses of panel data, and CRVEs are powerful tools for addressing serial correlation. For that reason, it seems natural to check test performance when ϵ_{ig} is serially correlated as an AR(1) process. [MacKinnon and Webb \(2018\)](#) find that a simple analytic test using a CRVE is less reliable when the errors are heteroskedastic. Following that paper as well as several other simulation studies of cluster-robust inference ([Cameron, Gelbach and Miller, 2008](#); [Djogbenou, MacK-](#)

innon and Nielsen, 2019), I also test robustness to the error variance differing across clusters.


In Figure 5, I plot rejection rates when the errors have a normal distribution and when they have non-normal distributions. I showed above that the degree of a given test's over- or under-rejection often depends on the test size α , so I continue to give results for both 1% and 5% tests here. I selected distributions with substantially different third and fourth moments than the normal distribution. The Laplace distribution is leptokurtic, the uniform distribution is platykurtic, and the log-normal distribution is skewed right. When $J = 200$ and $\phi = 1$ (panels 5a and 5d), so that there are many effective clusters, the different error distributions do not matter and every test rejects at the correct rate because all of the tests (including my test) are asymptotically consistent. When $J = 5$ (panels 5b and 5e) and when $\phi = 20$ (panels 5c and 5f), my test is not quite exact for non-normal error distributions, but it performs better than any of the other tests.

Next, I check test performance when the errors are serially correlated. Figure 6 shows the results of simulations where the errors are distributed as a stationary AR(1) process: $\epsilon_{i,g} = \psi\epsilon_{i-1,g} + \sqrt{1-\psi^2}\epsilon_{i,g}^*$, where $\epsilon_{i,g}^* \sim N(0,1)$. It bears repeating that when $J = 200$ and $\phi = 1$ (panels 6a and 6d), every test rejects at the correct rate because all of the tests (including my test) are asymptotically consistent. When $J = 5$ (panels 6b and 6e) and when $\phi = 20$ (panels 6c and 6f), my test is still nearly exact.

Finally, in Figure 7, I show the rejection rates for heteroskedastic errors. Specifically, I use a variant of the error distribution used to explore heteroskedasticity in Djogbenou, MacKinnon and Nielsen (2019); I let $\epsilon_{ig} = (1 + \xi(x_{1ig}^*)^2)^{\frac{1}{2}}\epsilon_{ig}^*$, where $\epsilon_{ig}^* \sim N(0,1)$. The error variance is higher for observations with greater magnitudes of x_{1ig} . When $J = 5$ (panels 7b and 7e), my test overrejects when ξ is high. When $\phi = 20$ (panels 7c and 7f), my test underrejects when ξ is high. Even in these cases, my test performs about as well as any other test.

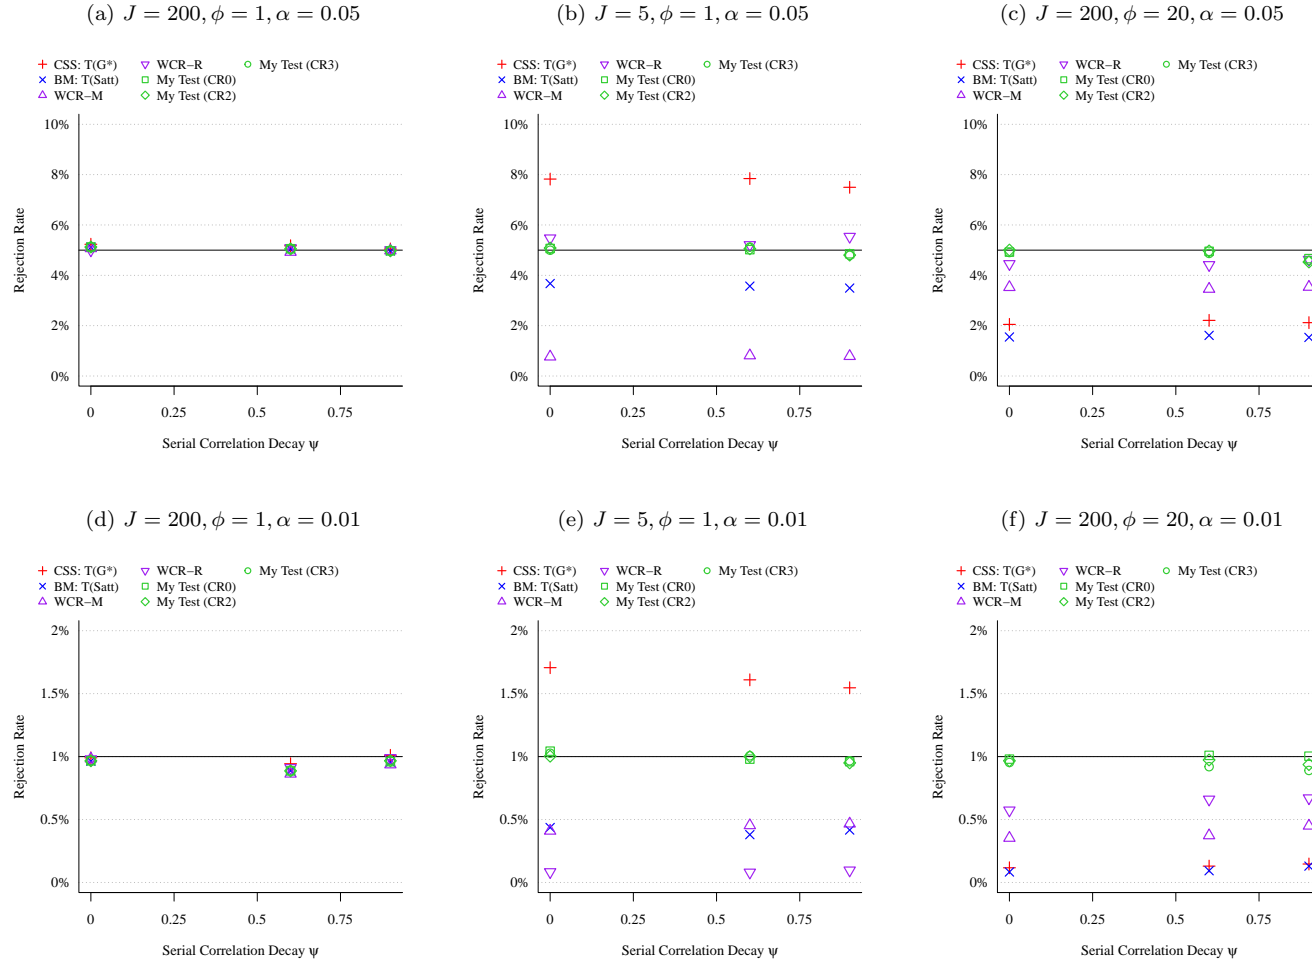
My test is only plausibly vulnerable to violations of Assumption 1 when the

Figure 5. : Comparison of Hypothesis Tests When the Error Distribution Varies



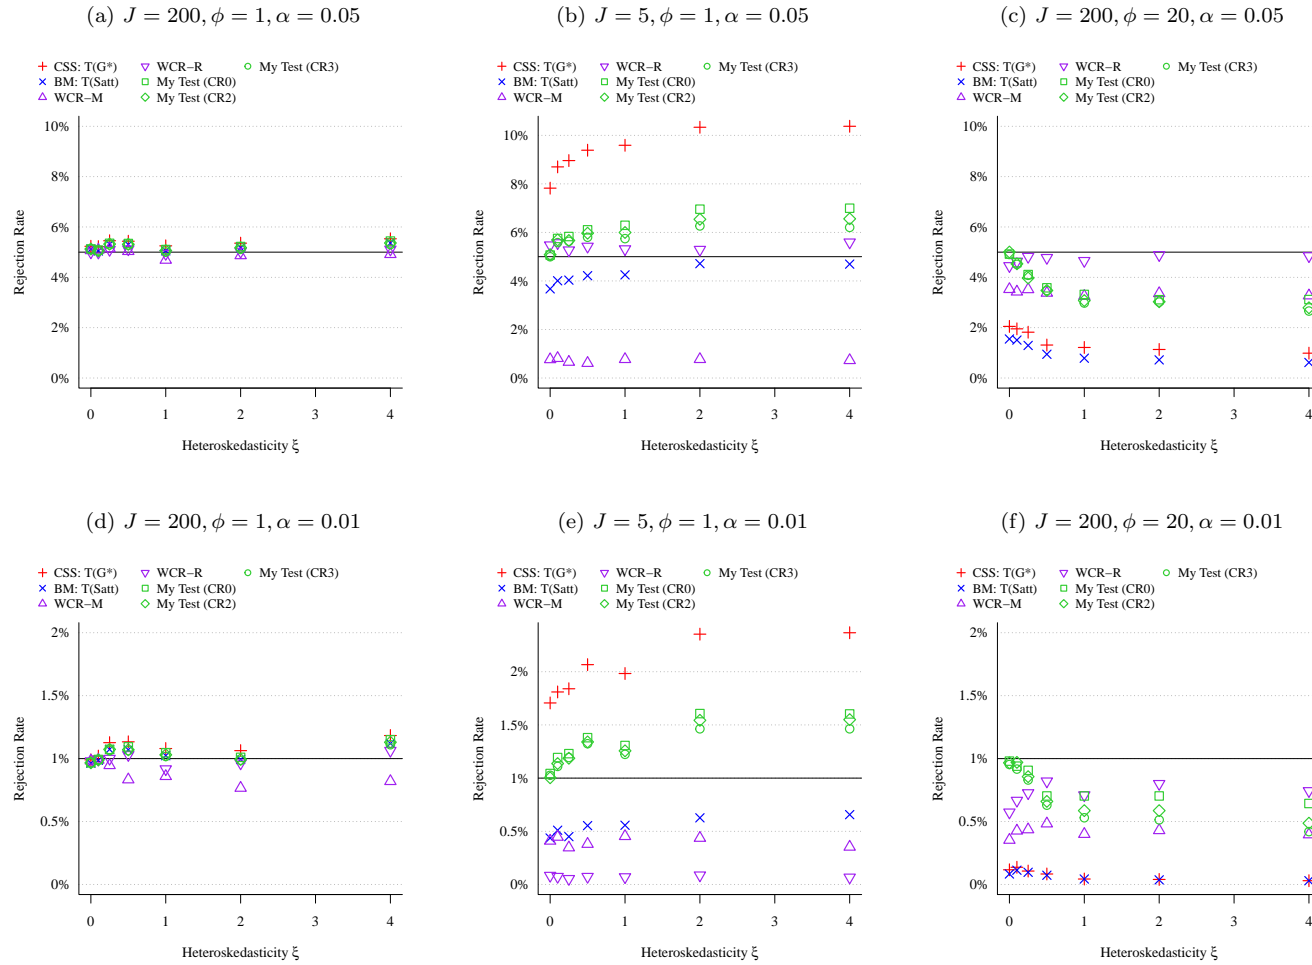
Notes: In these simulations, $G = 200$ and $N_1 = 5$. The errors are distributed as recentered and rescaled normal, Laplace, uniform, and log-normal distributions such that $\mathbb{E}(\epsilon_{ig}) = 0$ and $V(\epsilon_{ig}) = 1$. Each specification has 30,000 simulations.

Figure 6. : Comparison of Hypothesis Tests When the Errors are AR(1)



Notes: In these simulations, $G = 200$ and $N_1 = 5$. The errors are distributed as $\epsilon_{i,g} = \psi\epsilon_{i-1,g} + \sqrt{1 - \psi^2}\epsilon_{i,g}^*$, where $\epsilon_{i,g}^* \sim N(0, 1)$. Each specification has 30,000 simulations.

Figure 7. : Comparison of Hypothesis Tests When the Errors are Heteroskedastic



Notes: In these simulations, $G = 200$ and $N_1 = 5$. The errors are distributed as $\epsilon_{ig} = (1 + \xi(x_{1ig}^*)^2)^{\frac{1}{2}} \epsilon_{ig}^*$, where $\epsilon_{ig}^* \sim N(0, 1)$. Each specification has 30,000 simulations.

sample is not asymptotically large. I have shown in this section that several straightforward violations of Assumption 1 don't seem to affect the performance of my test very much. In particular, my test is less affected by these violations than the other tests are affected by few effective clusters.

5. Empirical Application

In this section, I apply my test to the empirical setting in [Abouk and Adams \(2013\)](#). First, I estimate their model and demonstrate some of the features of the sample that make the differences between methods of hypothesis testing meaningful. After that, I perform a simulation exercise using their data to discern how well different tests work in a context like this one.

[Abouk and Adams \(2013\)](#) estimate the impact on fatal car crashes of state bans on texting-while-driving. They use crash-level data from the Fatality Analysis Reporting System (FARS) of the National Highway Traffic Safety Administration. They focus on single-vehicle, single-occupant crashes in the US between 2007 and 2010, and they aggregate to the level of the state-by-month. Alaska is excluded from the sample because of missing data, for a total of 49 states (and thus 49 clusters).

As discussed in [Abouk and Adams \(2013\)](#), the enforcement mechanisms behind some of the bans were stronger than others, and they define two dummy variables $StrongBan_{s,t}$ and $WeakBan_{s,t}$ based on those two types of bans. If $StrongBan_{s,t} = 1$, then police officers could pull over any driver for texting-while-driving as a primary offense. This means that an officer with reasonable suspicion that a driver is texting can initiate a traffic stop on that basis alone, with no other traffic violation occurring. If $WeakBan_{s,t} = 1$, then texting-while-driving was either a secondary offense only (so that citations would be incidental to traffic stops for other primary offenses) or the ban only applied to a small subset of the population (e.g., drivers under 21 years old). By definition, no state has $StrongBan_{s,t} = WeakBan_{s,t} = 1$, and it happens that many states had

$$StrongBan_{s,t} = WeakBan_{s,t} = 0.$$

I replicate their estimation:

$$\begin{aligned} \log FatalCrashes_{s,t} = & \beta_{Strong} StrongBan_{s,t} + \beta_{Weak} WeakBan_{s,t} \\ & + w_{s,t}\delta + \gamma_s + \theta_t + \epsilon_{s,t} \end{aligned}$$

where $FatalCrashes_{s,t}$ is the number of fatal, single-vehicle, single-occupant car crashes in state s in year t , $w_{s,t}$ are a set of time-varying demographic and economic characteristics (e.g., unemployment rate), γ_s are state fixed effects, and θ_t are month-by-year fixed effects. The errors $\epsilon_{s,t}$ are clustered by state.

Table 1 shows coefficient estimates for $\hat{\beta}_{Strong}$ and $\hat{\beta}_{Weak}$. A “strong” ban reduced crashes by 8.1%. A “weak” ban is associated with an increase in crashes by 7.5%, but the estimate is imprecise.

I also give three different measures of the effective sample size. There were 49 states in the sample. I calculate the effective number of clusters G^* according to Carter, Schnepel and Steigerwald (2017) and the Satterthwaite-approximated degrees of freedom m from Bell and McCaffrey (2002). Formulas for both G^* and m can be found in Appendix C; they are functions of the covariate design matrix X , and in particular they are sensitive to the number of clusters, the number of clusters with treatment variation, cluster size heterogeneity, and treatment intensity heterogeneity. All three of these measures are used as the degrees of freedom v in different hypothesis tests.

One way to think about these measures is as measures of the variance of the test statistic t_0 . They represent the number of observations such that, if the sample was i.i.d., the variance of the t_0 would match the current sample. The measures differ between coefficients because test statistics based on different hypotheses have different behavior. The measures differ within a coefficient because they make different assumptions and simplifications ⁶.

⁶For a deeper discussion of the assumptions and simplifications behind the effective number of clusters G^* from Carter, Schnepel and Steigerwald (2017) and the Satterthwaite-approximated degrees

Table 1—: Impact of Texting-While-Driving Bans On Log(Number of Crashes)

	Independent Variable:	
	Strong Ban	Weak Ban
Coefficient Estimate	-0.081 (0.025)	0.075 (0.037)
# of Clusters (States)	49	49
G*, CSS Effective # of Clusters	11.6	7.2
m, BM-Satterthwaite DOF	10.7	5.7
P-Values for $H_0 : \beta = 0$		
No Adjustment	0.001	0.038
Hansen, T(G-1)	0.011	0.212
CSS, T(G*)	0.007	0.075
BM, T(m)	0.012	0.152
WCR-R Bootstrap	0.019	0.193
WCR-M Bootstrap	0.035	0.189
My Test (CR0)	0.008	0.111
My Test (CR2)	0.009	0.146
My Test (CR3)	0.011	0.190

Notes: Both columns contain estimates from the same regression. That regression also included state and year fixed effects as well as a set of time-varying state-level economic and demographic controls. Standard errors (in parentheses) are clustered by state. The effective number of clusters is calculated according to [Carter, Schnepel and Steigerwald \(2017\)](#) and the Satterthwaite-approximated degrees of freedom is calculated according to [Bell and McCaffrey \(2002\)](#). P-values are calculated for all of the hypothesis tests discussed in Section 3.

These measures suggest that the asymptotic properties of the standard CRVE might not dominate the behavior of a cluster-robust hypothesis test. Therefore, the differences between different methods of hypothesis testing might matter, especially for $\hat{\beta}_{Weak}$, where the Satterthwaite degrees of freedom is only 5.7.

In the bottom half of Table 1, I show the p-values given by each test discussed in Section 3. $\hat{\beta}_{Strong}$ is significantly different from 0 at the 5% level according to every test. $\hat{\beta}_{Weak}$ is significantly different from 0 at the 5% level according to the test that makes no adjustment to [White \(1984\)](#), using \hat{V}_{CR0} with critical values of freedom m from [Bell and McCaffrey \(2002\)](#), see Appendix C

drawn from the standard normal distribution, but it is not significant according to any other test. The adjustments to the CRVE and to critical value selection matter when the effective sample size is small.

It is not obvious how to interpret one p-value as more correct than another. A hypothesis test as a whole is exact if it rejects a test of size α with probability equal to α . Applying several tests in one sample – that is, in one realization of the error distribution – does not distinguish good tests from bad tests. However, I have performed a simulation exercise to learn which tests work well in samples like this one.





In this exercise, I randomize the values of the treatment variables $StrongBan_{s,t}$ and $WeakBan_{s,t}$. In each simulation, each state s is randomly assigned another state \tilde{s} (without replacement). Then I assign $RStrongBan_{s,t} = StrongBan_{\tilde{s},t}$, and I assign $RWeakBan_{s,t} = WeakBan_{\tilde{s},t}$, so that each state s is assigned the entire sequence of treatment variables from state \tilde{s} .

I have assigned these treatments randomly and after the fact, so it is safe to assume that $\beta_{RStrong} = \beta_{RWeak} = 0$. Since $H_0 : \beta_{RStrong} = 0$ and $H_0 : \beta_{RWeak} = 0$ are true hypotheses, an exact hypothesis test should reject with probability equal to the nominal size of the test. Figure 8 shows the rejection rates for each test method across all the simulations.

Hansen substantially overrejects $H_0 : \beta_{RWeak} = 0$ for a test size of $\alpha = 0.01$, while CSS substantially overrejects $H_0 : \beta_{RStrong} = 0$ for a test size of $\alpha = 0.01$. CSS also seems to overreject somewhat for a test size of $\alpha = 0.05$ on either coefficient.

The variant of my test that uses \hat{V}_{CR3} is about as close to exact as any other test for both coefficients and test sizes. Also, it's worth noting that in Figure 8 the rejection rates of all variants of my test tend to be between those of the other tests. This comports with the Monte Carlo results in Section 4, where some tests over-reject and some under-reject.

The important features of this empirical setting are not unique. Staggered

Figure 8. : Simulation Exercise Based On [Abouk and Adams \(2013\)](#)(a) $H_0 : \beta_{RStrong} = 0, \quad \alpha = 0.05$ (b) $H_0 : \beta_{RStrong} = 0, \quad \alpha = 0.01$ (c) $H_0 : \beta_{RWeak} = 0, \quad \alpha = 0.05$ (d) $H_0 : \beta_{RWeak} = 0, \quad \alpha = 0.01$ 

Notes: These plots show the rejection rate of each test when the values of the treatment variables $StrongBan_{s,t}$ and $WeakBan_{s,t}$ have been randomized over 30,000 simulations.

changes in public policy across states or counties provide useful natural experiments for learning about the impact of public policy. A good hypothesis test is necessary for understanding how much can be learned from a given experiment.

6. Conclusion

In this paper, I have proposed a hypothesis test for inference in clustered samples with cluster-level fixed effects. My test is robust to samples with few clusters, few clusters with treatment variation, cluster size outliers, or treatment

intensity outliers. In addition to being consistent in large samples, my test is also exact under normal, homoskedastic errors, and in simulations it performs well compared to other methods from the literature under other assumptions about the errors.

Many samples are large enough that the conventional cluster-robust test will work fine, drawing critical values from the standard normal distribution. However, some samples are not as large as they seem, in the sense of the distribution of the test statistic.

Samples with many observations but few clusters require adjustment to the conventional test. Samples with many clusters but only a few with treatment variation require adjustment. And samples where the residual treatment variation (conditional on covariates) is concentrated in a small number of clusters require adjustment. Covariates can hide the features that change the behavior of the test statistic. It is worthwhile to use a test that is simply robust to these features.

REFERENCES

- Abouk, Rahi, and Scott Adams.** 2013. “Texting Bans and Fatal Accidents on Roadways: Do They Work? Or Do Drivers Just React to Announcements of Bans?” *American Economic Journal: Applied Economics*, 5(2): 179–199.
- Arellano, M.** 1987. “Computing Robust Standard Errors for Within-groups Estimators.” *Oxford Bulletin of Economics and Statistics*, 49(4): 431–434.
- Bell, Robert M., and Daniel F. McCaffrey.** 2002. “Bias reduction in standard errors for linear regression with multi-stage samples.” *Survey Methodology*, 28(2): 169–182. ISBN: 0714-0045.
- Bertrand, M., E. Duflo, and S. Mullainathan.** 2004. “How Much Should We Trust Differences-In-Differences Estimates?” *The Quarterly Journal of Economics*, 119(1): 249–275.

- Bound, John, Breno Braga, Gaurav Khanna, and Sarah Turner.** 2020. "A Passage to America: University Funding and International Students." *American Economic Journal: Economic Policy*, 12(1): 97–126.
- Cameron, A. Colin, and Douglas L. Miller.** 2015. "A Practitioner's Guide to Cluster-Robust Inference." *Journal of Human Resources*, 50(2): 317–372.
- Cameron, A. Colin, Jonah B. Gelbach, and Douglas L. Miller.** 2008. "Bootstrap-Based Improvements for Inference with Clustered Errors." *Review of Economics and Statistics*, 90(3): 414–427.
- Carter, Andrew V., Kevin T. Schnepel, and Douglas G. Steigerwald.** 2017. "Asymptotic Behavior of a t -Test Robust to Cluster Heterogeneity." *The Review of Economics and Statistics*, 99(4): 698–709.
- Djogbenou, Antoine A., James G. MacKinnon, and Morten Ørregaard Nielsen.** 2019. "Asymptotic theory and wild bootstrap inference with clustered errors." *Journal of Econometrics*, 212(2): 393–412.
- Hansen, Bruce E.** 2021. *Econometrics*. Princeton, New Jersey: Princeton University Press.
- Hansen, Christian B.** 2007. "Asymptotic properties of a robust variance matrix estimator for panel data when is large." *Journal of Econometrics*, 141(2): 597–620.
- Karaivanov, Alexander, Shih En Lu, Hitoshi Shigeoka, Cong Chen, and Stephanie Pamplona.** 2021. "Face masks, public policies and slowing the spread of COVID-19: Evidence from Canada." *Journal of Health Economics*, 78: 102475.
- Kuziemko, Ilyana, Katherine Meckel, and Maya Rossin-Slater.** 2018. "Does Managed Care Widen Infant Health Disparities? Evidence from Texas Medicaid." *American Economic Journal: Economic Policy*, 10(3): 255–283.

- Liang, Kung-Yee, and Scott L. Zeger.** 1986. “Longitudinal data analysis using generalized linear models.” *Biometrika*, 73(1): 13–22.
- MacKinnon, James G., and Matthew D. Webb.** 2018. “The wild bootstrap for few (treated) clusters.” *The Econometrics Journal*, 21(2): 114–135.
- Myers, Caitlin Knowles.** 2017. “The Power of Abortion Policy: Reexamining the Effects of Young Women’s Access to Reproductive Control.” *Journal of Political Economy*, 125(6): 2178–2224.
- White, Halbert.** 1984. *Asymptotic Theory for Econometricians. Economic Theory, Econometrics, and Mathematical Economics*, Orlando, Florida:Academic Press.

APPENDIX A Proof of Theorem 1

The null hypothesis H_0 is true if $c_0^T \beta = a_0$. Assumption 1 holds when $\epsilon_g \sim N(0, \sigma^2 \Omega_g)$, where $\Omega_g = (1 - \rho)I_g + \rho \iota_g \iota_g^T$ (σ and ρ may be unknown).

Theorem 1 says that when H_0 is true and Assumption 1 holds, the CDF of the squared test statistic is a known function of the design matrix X and the hypothesis H_0 :

$$P(t_0^2 < q \mid X) = L(q; X, H_0)$$

In this section, I will prove Theorem 1 by deriving $L(\cdot; \cdot, \cdot)$. I give some notation to help with the analysis, I prove that $L(\cdot; \cdot, \cdot)$ is known in principle, and I show how $L(\cdot; \cdot, \cdot)$ can be calculated quickly in practice.

A.1 Notation

Let N_g be the number of observations in cluster g , and let $N = \sum_g N_g$. Then let I_g be an identity matrix of size N_g , and let ι_g be a column vector of length N_g whose elements are all 1. So then let $M_g = I_g - \frac{1}{N_g} \iota_g \iota_g^T$, and note that $M_g \iota_g = 0$.

Recall that:

$$y_{ig} = x_{ig} \beta + \gamma_g + \epsilon_{ig}$$

For identification, we have that $\mathbb{E}(\epsilon_{ig} \mid x_{ig}) = 0$, and for inference, we have that $\mathbb{E}(\epsilon_{ig} \epsilon_{jg'}) = 0$. Y_g is $(N_g \times 1)$, stacking up the dependent variable within cluster g , and X_g is $(N_g \times K)$, stacking up the covariates within cluster g . Then $\epsilon_g = Y_g - X_g \beta - \gamma_g \iota_g$.

For fixed effects absorption – that is, to absorb γ_g – the cluster-level means of y_{ig} and x_{ig} are subtracted from the individual-level y_{ig} and x_{ig} , respectively. Stated another way:

$$\ddot{Y}_g = M_g Y_g$$

$$\ddot{X}_g = M_g X_g$$

$$\ddot{\epsilon}_g = M_g \epsilon_g = M_g(Y_g - X_g - \gamma_g \iota_g) = \ddot{Y}_g - \ddot{X}_g \beta$$

Also recall that $\hat{\beta}$ is the fixed effects estimator of β , and $\hat{V}(\hat{\beta})$ is the CRVE:

$$\begin{aligned}\hat{\beta} &= (\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T \ddot{Y} \\ \hat{\epsilon}_g &= \ddot{Y}_g - \ddot{X}_g \hat{\beta} \\ \hat{V}(\hat{\beta}) &= (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T A_g \hat{\epsilon}_g \hat{\epsilon}_g^T A_g^T \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1}\end{aligned}$$

where (A_g) are adjustment matrices according to Bell and McCaffrey (2002); for \hat{V}_{CR0} , $A_g = I_g$.

Finally, recall the hypothesis H_0 and the test statistic t_0 :

$$\begin{aligned}H_0 : c_0^T \beta &= a_0 \\ t_0 &= \frac{c_0^T \hat{\beta} - a_0}{\sqrt{c_0^T \hat{V}(\hat{\beta}) c_0}}\end{aligned}$$

A.2 Main Proof

First, I show that the CDF of t_0^2 at a particular quantile q can be written as the CDF at 0 of a linear combination of χ^2 random variables. Second, I show that it is possible to determine the coefficients of that linear combination. Third, I give a formula for $L(q; X, H_0)$.

Since the hypothesis H_0 holds:

$$\begin{aligned}t_0^2 &= \frac{(c_0^T \hat{\beta} - a_0)^2}{c_0^T \hat{V}(\hat{\beta}) c_0} \\ &= \frac{(c_0^T (\hat{\beta} - \beta))^2}{c_0^T \hat{V}(\hat{\beta}) c_0} \\ &= \frac{c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T \ddot{\epsilon} \ddot{\epsilon}^T \ddot{X} (\ddot{X}^T \ddot{X})^{-1} c_0}{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T A_g \hat{\epsilon}_g \hat{\epsilon}_g^T A_g^T \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0}\end{aligned}$$

Let $H = \ddot{X}(\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T$, let I be an identity matrix of size N , and let $(I - H)_g$ be the rows of $(I - H)$ corresponding to cluster g . Then using the fact that $\hat{\epsilon}_g = \ddot{\epsilon}_g - \ddot{X}_g(\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T \ddot{\epsilon} = (I - H)_g \ddot{\epsilon}$, I continue:

$$\begin{aligned} t_0^2 &= \frac{c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T \ddot{\epsilon} \ddot{\epsilon}^T \ddot{X} (\ddot{X}^T \ddot{X})^{-1} c_0}{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T A_g (I - H)_g \ddot{\epsilon} \ddot{\epsilon}^T (I - H)_g^T A_g^T \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0} \\ &= \frac{\ddot{\epsilon}^T \ddot{X} (\ddot{X}^T \ddot{X})^{-1} c_0 c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T \ddot{\epsilon}}{\ddot{\epsilon}^T \left(\sum_g (I - H)_g^T A_g^T \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0 c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T A_g (I - H)_g \right) \ddot{\epsilon}} \end{aligned}$$

Now let $d_0 = \ddot{X}(\ddot{X}^T \ddot{X})^{-1} c_0$, and for $g \geq 1$, let $d_g = (I - H)_g^T A_g^T \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0$.

Then it follows that:

$$\begin{aligned} t_0^2 &= \frac{\ddot{\epsilon}^T d_0 d_0^T \ddot{\epsilon}}{\ddot{\epsilon}^T \left(\sum_g d_g d_g^T \right) \ddot{\epsilon}} \\ P(t_0^2 < q \mid X) &= P \left(\frac{\ddot{\epsilon}^T d_0 d_0^T \ddot{\epsilon}}{\ddot{\epsilon}^T \left(\sum_g d_g d_g^T \right) \ddot{\epsilon}} < q \mid X \right) \\ &= P \left(\frac{1}{q} \ddot{\epsilon}^T (d_0 d_0^T) \ddot{\epsilon} - \ddot{\epsilon}^T \left(\sum_g d_g d_g^T \right) \ddot{\epsilon} < 0 \mid X \right) \\ &= P \left(\ddot{\epsilon}^T \left(\frac{1}{q} d_0 d_0^T - \sum_g d_g d_g^T \right) \ddot{\epsilon} < 0 \mid X \right) \end{aligned}$$

Let $D_+ = [d_0 \ d_1 \dots d_g \dots d_G]$ and $D_- = [\frac{1}{q} d_0 \ -d_1 \dots -d_g \dots -d_G]$, so that:

$$(5) \quad P(t_0^2 < q \mid X) = P \left(\ddot{\epsilon}^T (D_+ D_-^T) \ddot{\epsilon} < 0 \mid X \right)$$

Now using Assumption 1, the errors before fixed effects absorption are distributed $\epsilon_g \sim N(0, \sigma^2 \Omega_g)$. Then the errors after absorption, $\ddot{\epsilon}_g = M_g \epsilon_g$, are distributed $\ddot{\epsilon}_g \sim N(0, \sigma^2 \ddot{\Omega}_g)$, where:

$$\ddot{\Omega}_g = M_g \Omega_g M_g$$

$$\begin{aligned}
&= \left(I_g - \frac{1}{N_g} \iota_g \iota_g^T \right) ((1 - \rho) I_g + \rho \iota_g \iota_g^T) M_g \\
&= \left(I_g - \frac{1}{N_g} \iota_g \iota_g^T \right) ((1 - \rho) I_g) M_g \\
&= (1 - \rho) \left(I_g - \frac{1}{N_g} \iota_g \iota_g^T \right) M_g \\
\ddot{\Omega}_g &= (1 - \rho) M_g M_g
\end{aligned}$$

Then let M be an $(N \times N)$ block-diagonal matrix where the g -th block is M_g . And let η be joint normal with $\eta \sim N(0, I)$, such that $\ddot{\epsilon} = \sigma(1 - \rho)^{\frac{1}{2}} M \eta$. Substituting this into (5):

$$\begin{aligned}
L(q; X, H_0) &= P(t_0^2 < q \mid X) \\
&= P\left(\eta^T M \sigma(1 - \rho)^{\frac{1}{2}} (D_+ D_-^T) \sigma(1 - \rho)^{\frac{1}{2}} M \eta < 0 \mid X\right) \\
&= P\left(\eta^T M D_+ D_-^T M \eta < 0 \mid X\right)
\end{aligned}$$

At this point, the right-hand side is a function only of known quantities and random variables with known distributions. Neither σ nor ρ appear; the behavior of the test statistic t_0 does not depend on them.

A.3 Calculating $L(q; X, H_0)$ in Practice

I have shown that $L(q; X, H_0) = P(\eta^T M D_+ D_-^T M \eta < 0 \mid X)$. In this section, I will explain a method for calculating $L(q; X, H_0)$ quickly in practice.

Let $Q_{(N \times N)} = M D_+ D_-^T M$. Then let S be the orthogonal⁷ matrix of eigenvectors of Q , let λ^* be an $(N \times 1)$ column vector whose elements are the eigenvalues of Q , and let Λ be an $(N \times N)$ diagonal matrix whose diagonal elements are also the eigenvalues of Q . Note that since S is orthogonal, $S \eta \sim \eta$. Then:

$$P(t_0^2 < q \mid X) = P(\eta^T S \Lambda S^T \eta < 0 \mid X)$$

⁷An orthogonal S can always be found because $Q = M \left(\frac{1}{q} d_0 d_0^T - \sum_g d_g d_g^T \right) M$ is symmetric.

$$= P(\eta^T \Lambda \eta < 0 \mid X)$$

Let w be an N -vector of independent random variables such that $\forall i, w_i \sim \chi_1^2$. Then:

$$(6) \quad P(t_0^2 < q \mid X) = P(w^T \lambda^* < 0 \mid X)$$

Thus, I have shown that the CDF of t_0^2 at q can be written as the CDF at 0 of a linear combination of independent χ_1^2 random variables. Next, I find the non-zero elements of λ^* ; it will be the case that λ^* has no more than $G + 1$ non-zero elements.

In principle, the vector of eigenvalues λ^* can be found by eigendecomposing Q . However, since Q is $(N \times N)$, that might be inconvenient in practice. Instead, it is sufficient to find the non-zero eigenvalues of $D_-^T M D_+$, which are the same as the non-zero eigenvalues of $Q = M D_+ D_-^T M$.

To see why, suppose that λ_j is a non-zero eigenvalue of $M D_+ D_-^T M$ corresponding to the eigenvector s_j . Note that M is idempotent. Therefore:

$$\begin{aligned} M D_+ D_-^T M s_j &= \lambda_j s_j \\ D_-^T M M D_+ D_-^T M s_j &= D_-^T M \lambda_j s_j \\ D_-^T M D_+ (D_-^T M s_j) &= \lambda_j (D_-^T M s_j) \end{aligned}$$

Thus, λ_j is an eigenvalue of $D_-^T M D_+$ corresponding to the eigenvector $D_-^T M s_j$. And since $D_-^T M D_+$ is a $(G + 1 \times G + 1)$ matrix, it has no more than $G + 1$ non-zero eigenvalues. Letting λ be a $(G + 1 \times 1)$ vector whose elements are the eigenvalues of $D_-^T M D_+$, and in an abuse of notation letting w now be $(G + 1 \times 1)$, we have that:

$$(7) \quad P(t_0^2 < q \mid X) = P(w^T \lambda < 0 \mid X)$$

The CDF of a linear combination of independent χ_1^2 random variables $w^T \lambda$ is given by Imhof (1961)⁸:

$$P(w^T \lambda < 0 \mid X) = \frac{1}{2} - \frac{1}{\pi} \int_0^\infty \frac{\sin\left(\frac{1}{2} \sum_{j=1}^{G+1} \tan^{-1}(\lambda_j u)\right)}{u \prod_{j=1}^{G+1} (1 + \lambda_j^2 u^2)^{\frac{1}{4}}} du$$

So the CDF of t_0^2 at q can be written as the CDF at 0 of a linear combination of $G + 1$ independent χ_1^2 random variables, and it is possible to calculate the coefficients λ as a function of the design matrix X , the hypothesis H_0 , and the quantile q . In summary:

$$d_0 = \ddot{X}(\ddot{X}^T \ddot{X})^{-1} c_0, \quad d_g = (I - H)_g^T A_g^T \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0,$$

$$D_+ = [d_0 \quad d_1 \dots d_g \dots d_G], \quad D_- = \begin{bmatrix} \frac{1}{q} d_0 & -d_1 \dots -d_g \dots -d_G \end{bmatrix},$$

M block-diagonal, with g -th block M_g ,

λ are the eigenvalues of $D_-^T M D_+$, and

$$L(q; X, H_0) = P(t_0^2 < q \mid X)$$

$$= \frac{1}{2} - \frac{1}{\pi} \int_0^\infty \frac{\sin\left(\frac{1}{2} \sum_{j=1}^{G+1} \tan^{-1}(\lambda_j u)\right)}{u \prod_{j=1}^{G+1} (1 + \lambda_j^2 u^2)^{\frac{1}{4}}} du$$

And this is what I set out to find.

My test involves selecting a critical value:

$$q^*(\alpha; X, H_0) = \sqrt{L^{-1}(1 - \alpha; X, H_0)}$$

And I reject H_0 if $|t_0| > q^*$. Thus, under Assumption 1, my test is exact, so that the rate at which a true hypothesis is rejected is equal to the nominal size of the

⁸This can be calculated quickly by numerical integration, with a high degree of precision, using the *imhof()* function from the R package *CompQuadForm*.

test:

$$\begin{aligned} P(|t_0| > q^*) &= P(t_0^2 > (q^*)^2) = 1 - L((q^*)^2; X, H_0) \\ &= \alpha \end{aligned}$$

APPENDIX B Asymptotic Consistency

In this section, I show that my test is asymptotically consistent (without Assumption 1). This result draws on the first theorem from [Carter, Schnepel and Steigerwald \(2017\)](#).

The null hypothesis H_0 is true if $c_0^T \beta = a_0$. Assumption 2 requires that the errors have fourth moments, and Assumption 3 requires that the observations aren't too concentrated in a small number of clusters.

Let $\alpha^* = P(|t_0| > q^*(\alpha; X, H_0))$ be the rejection rate of my test – the rate at which my test rejects a true hypothesis H_0 .

Theorem 2 says that when H_0 is true and Assumptions 2 and 3 hold, the rejection rate of my test converges to the nominal size of the test:

$$\alpha^* \xrightarrow{p} \alpha$$

Recall that $\mathbb{E}(\epsilon_{ig}\epsilon_{jg'}) = 0$, so that the errors are uncorrelated across clusters. According to the first theorem in [Carter, Schnepel and Steigerwald \(2017\)](#), it follows that $t_0 \xrightarrow{d} N(0, 1)$.

Now, consider a counterfactual data generating process:

$$\tilde{y}_{ig} = x_{ig}\beta + \gamma_g + \tilde{\epsilon}_{ig}$$

where $\tilde{\epsilon} \sim N(0, I)$, so that the counterfactual errors are normal, i.i.d., and homoskedastic. Let \tilde{t}_0 be the test statistic that would be generated for H_0 using \hat{V}_{CR0} in the counterfactual data generating process where the errors are $\tilde{\epsilon}$ rather than ϵ :

$$\begin{aligned} \tilde{\beta} &= \beta + (\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T M \tilde{\epsilon} \\ \hat{\tilde{\epsilon}} &= \ddot{Y} - \ddot{X} \tilde{\beta} \end{aligned}$$

$$\tilde{V} = (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T \hat{\epsilon}_g \hat{\epsilon}_g^T \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1}$$

$$\tilde{t}_0 = \frac{c_0^T (\tilde{\beta} - \beta)}{\sqrt{c_0^T \tilde{V} c_0}}$$

Since $\tilde{\epsilon}$ meets the conditions of Assumptions 2 and 3, the first theorem in [Carter, Schnepel and Steigerwald \(2017\)](#) applies, so $\tilde{t}_0 \xrightarrow{P} N(0, 1)$. And whereas $\tilde{\epsilon}$ also meets the conditions of Assumption 1, it is also the case that Theorem 1 applies, so that $P(|\tilde{t}_0| \geq q^*(\alpha; X, H_0) \mid X) = \alpha$.

Note that for the case of stochastic X with PDF $f(\cdot)$, it follows from Theorem 1 and the law of total probability that:

$$\begin{aligned} P(|\tilde{t}_0| \geq q^*(\alpha; X, H_0)) &= \int P(|\tilde{t}_0| \geq q^* \mid X) f(X) dX \\ &= \int \alpha f(X) dX \\ &= \alpha \end{aligned}$$

Therefore:

$$\begin{aligned} \alpha^* - \alpha &= P(|t_0| \geq q^*) - P(|\tilde{t}_0| \geq q^*) \\ &= (1 - P(t_0 < q^*) + P(t_0 < -q^*)) \\ &\quad - (1 - P(\tilde{t}_0 < q^*) + P(\tilde{t}_0 < -q^*)) \\ &\xrightarrow{P} (\Phi(-q^*) - \Phi(q^*)) - (\Phi(-q^*) - \Phi(q^*)) \\ \alpha^* - \alpha &\xrightarrow{P} 0 \\ \alpha^* &\xrightarrow{P} \alpha \end{aligned}$$

where $\Phi(\cdot)$ is the CDF of the standard normal distribution. In summary:

- 1) The true test statistic t_0 and the counterfactual test statistic \tilde{t}_0 converge to the same distribution.

- 2) My test is exact for the counterfactual DGP.
- 3) My test converges to an exact test for any DGP meeting Assumptions 2 and 3.

APPENDIX C Approximating t_0 as $T(v)$

In this section, I will discuss how previous tests have selected critical values for the test statistic t_0 by approximating its distribution as $T(v)$, a t -distribution with v degrees of freedom. [Bell and McCaffrey \(2002\)](#) use $T(m)$, where m is the degrees of freedom from a Satterthwaite approximation, and [Carter, Schnepel and Steigerwald \(2017\)](#) use $T(G^*)$, where G^* is the effective number of clusters. The purpose of this section will be to make sense of the assumptions and simplifications that are necessary to rationalize those tests in a finite sample.

Consider the definition of a t -distributed random variable with v degrees of freedom as τ :

$$\begin{aligned}
 \tau &\sim T(v) \\
 \tau &= \frac{z}{\sqrt{\Upsilon}}, \quad \text{where } z \sim N(0, 1), \\
 v\Upsilon &\sim \chi_v^2, \text{ and} \\
 \Upsilon &\perp z
 \end{aligned}
 \tag{8}$$

The intuition behind the approximations below is that, under Assumption 1, the test statistic resembles this structure superficially; t_0 is a ratio of a normal random variable divided by the square root of a sum of squared normals. For simplicity, assume a non-stochastic covariate design matrix X . Let $z_0 = \frac{c_0^T(\hat{\beta} - \beta)}{\sqrt{c_0^T V(\hat{\beta}) c_0}}$ and let $\Upsilon_0 = \frac{c_0^T \hat{V}(\hat{\beta}) c_0}{c_0^T V(\hat{\beta}) c_0}$. Now consider the test statistic for a true hypothesis $H_0 : c_0^T t_0 = a_0$:

$$\begin{aligned}
 t_0 &= \frac{c_0^T \hat{\beta} - a_0}{\sqrt{c_0^T \hat{V}(\hat{\beta}) c_0}} \\
 &= \frac{c_0^T (\hat{\beta} - \beta)}{\sqrt{c_0^T \hat{V}(\hat{\beta}) c_0}}
 \end{aligned}$$

$$\begin{aligned}
&= \frac{c_0^T(\hat{\beta} - \beta)}{\sqrt{c_0^T V(\hat{\beta}) c_0}} \times \sqrt{\frac{c_0^T V(\hat{\beta}) c_0}{c_0^T \hat{V}(\hat{\beta}) c_0}} \\
&= \frac{z_0}{\sqrt{\Upsilon_0}}
\end{aligned}$$

If it were the case that $z_0 \sim N(0, 1)$, that $(v_0 \Upsilon_0) \sim \chi_{v_0}^2$ for some v_0 , and that $z_0 \perp \Upsilon_0$, then t_0 would in fact have a t -distribution with v_0 degrees of freedom.

C.1 Approximation in *Bell and McCaffrey (2002)*

In this section, I will show how the approximation to a t -distribution is constructed in *Bell and McCaffrey (2002)*. As in Appendix A, I define N_g as the number of observations in cluster g , I_g as an identity matrix of size N_g , and ι_g as a column vector of length N_g whose elements are all 1. Additionally, $M_g = I_g - \frac{1}{N_g} \iota_g \iota_g^T$, and M is a block-diagonal matrix whose g -th block is M_g . Also, recall that $\ddot{\epsilon}_g \sim N(0, \sigma^2(1 - \rho)M_g)$. Furthermore, define:

$$\begin{aligned}
d_g &= (I - H)_g^T A_g^T \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0, \\
D_{(N \times G)} &= [d_1 \dots d_g \dots d_G] \\
\lambda_{(G \times 1)}^{BM} &\text{ are the eigenvalues of } D^T M D \\
m &= \frac{\left(\sum_g \lambda_g^{BM}\right)^2}{\sum_g (\lambda_g^{BM})^2}
\end{aligned}$$

Bell and McCaffrey (2002) calculate a test statistic using the variance estimator \hat{V}_{CR2} (so $A_g = (I_g - H_{gg})^{\frac{1}{2}}$), and they select critical values from $T(m)$. They assume that $\epsilon \sim N(0, \sigma^2 I_N)$; since Assumption 1 is sufficient (and weaker), I will refer to that assumption. Assumption 1 implies that:

- 1) \hat{V}_{CR2} is unbiased: $\mathbb{E}(\hat{V}_{CR2}) = V(\hat{\beta})$
- 2) $\hat{\epsilon} \perp \hat{\beta}$

It follows immediately that $z_0 \perp \Upsilon_0$. Furthermore, observe that under Assumption 1, $z_0 \sim N(0, 1)$. This follows from the normality of the errors; $c_0^T \hat{\beta}$ is normal with mean $c_0^T \beta$ and variance $c_0^T V(\hat{\beta}) c_0$. Now following a similar logic as in Appendix A:

$$\begin{aligned}
c_0^T \hat{V}(\hat{\beta}) c_0 &= c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T A_g \hat{\epsilon}_g \hat{\epsilon}_g^T A_g \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0 \\
&= c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T A_g (I - H)_g \ddot{\epsilon} \ddot{\epsilon}^T (I - H)_g^T A_g \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0 \\
&= \sum_g \ddot{\epsilon}^T (I - H)_g^T A_g \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0 c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T A_g (I - H)_g \ddot{\epsilon} \\
&= \ddot{\epsilon}^T \left(\sum_g d_g d_g^T \right) \ddot{\epsilon}
\end{aligned}$$

Again recalling from Appendix A that $\ddot{\epsilon} \sim N(0, \sigma^2(1 - \rho)M)$, let $\eta = N(0, I)$ such that $\ddot{\epsilon} = \sigma(1 - \rho)^{\frac{1}{2}} M \eta$ and let w be a G -vector of independent χ_1^2 random variables. Also, the non-zero eigenvalues of $MDD^T M$ are the same as the non-zero eigenvalues of $D^T M D$. Then:

$$\begin{aligned}
c_0^T \hat{V}(\hat{\beta}) c_0 &= \sigma^2(1 - \rho) \eta^T M D D^T M \eta \\
&\sim \sigma^2(1 - \rho) w^T \lambda^{BM}
\end{aligned}$$

Now since the variance estimator is unbiased:

$$\begin{aligned}
c_0^T V(\hat{\beta}) c_0 &= \mathbb{E} \left(c_0^T \hat{V}(\hat{\beta}) c_0 \right) \\
&= \mathbb{E} \left(\sigma^2(1 - \rho) w^T \lambda^{BM} \right) \\
&= \sigma^2(1 - \rho) \sum_g \mathbb{E}(w_g) \lambda_g^{BM} \\
&= \sigma^2(1 - \rho) \sum_g \lambda_g^{BM}
\end{aligned}$$

It follows that:

$$\begin{aligned}\Upsilon_0 &= \frac{c_0^T \hat{V}(\hat{\beta}) c_0}{c_0^T V(\hat{\beta}) c_0} = \frac{\sum_g w_g \lambda_g^{BM}}{\sum_g \lambda_g^{BM}} \\ \mathbb{E}(\Upsilon_0) &= 1 \\ V(\Upsilon_0) &= \sum_g V(w_g) \left(\frac{\lambda_g^{BM}}{\sum_{g'} \lambda_{g'}^{BM}} \right)^2 = \frac{2}{m}\end{aligned}$$

Now, [Bell and McCaffrey \(2002\)](#) can apply the Satterthwaite approximation of $\Upsilon_0 \approx \Upsilon^{BM}$, where $m\Upsilon^{BM} \sim \chi_m^2$. The first two moments match, so that $\mathbb{E}(\Upsilon^{BM}) = 1$ and $V(\Upsilon^{BM}) = \frac{2}{m}$. Let $\tau^{BM} = \frac{z_0}{\sqrt{\Upsilon^{BM}}}$ be the approximated test statistic with Υ^{BM} substituted for Υ_0 . Then it is in fact the case that $\tau^{BM} \sim T(m)$.

In summary, [Bell and McCaffrey \(2002\)](#) make (a stronger version of) Assumption 1. Then they apply a Satterthwaite approximation to the denominator of the test statistic. As a result, they find that the approximated test statistic t^{BM} has a t distribution with m degrees of freedom, where $m = \frac{(\sum_g \lambda_g^{BM})^2}{\sum_g (\lambda_g^{BM})^2}$.

C.2 Approximation in [Carter, Schnepel and Steigerwald \(2017\)](#)

I will now show how [Carter, Schnepel and Steigerwald \(2017\)](#) approximate the test statistic t_0 as a t -distribution. Their method also involves a Satterthwaite approximation. They assume that $\forall g \epsilon_g \sim N(0, \sigma^2 \iota_g \iota_g^T)$; since Assumption 1 is once again sufficient (and weaker), I will refer to that assumption.

In [Carter, Schnepel and Steigerwald \(2017\)](#), the test statistic is calculated using the variance estimator \hat{V}_{CR0} (so $A_g = I_g$). Then, critical values are selected from $T(G^*)$, where G^* is what they refer to as the “effective number of clusters”:

$$\begin{aligned}\lambda_g^{CSS} &= c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T \mathbb{E}(\ddot{\epsilon}_g \ddot{\epsilon}_g^T) \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0 \\ &= \sigma(1 - \rho) c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T M_g \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0\end{aligned}$$

$$G^* = \frac{(\sum_g \lambda_g^{CSS})^2}{\sum_g (\lambda_g^{CSS})^2}$$

Recall that $\mathbb{E}(\ddot{\epsilon}_g \ddot{\epsilon}_g^T) = \sigma^2(1 - \rho)M_g$. Now returning to the test statistic, $t_0 = \frac{z_0}{\sqrt{\Upsilon_0}}$ and we can decompose Υ_0 into two parts:

$$\begin{aligned} \Upsilon_0 &= \frac{c_0^T \hat{V}(\hat{\beta}) c_0}{c_0^T V(\hat{\beta}) c_0} \\ &= \frac{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T A_g \hat{\epsilon}_g \hat{\epsilon}_g^T A_g \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0}{c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}^T \mathbb{E}(\ddot{\epsilon} \ddot{\epsilon}^T) \ddot{X} (\ddot{X}^T \ddot{X})^{-1} c_0} \\ &= \frac{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T (I - H)_g \ddot{\epsilon} \ddot{\epsilon}^T (I - H)_g^T \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0}{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T \mathbb{E}(\ddot{\epsilon}_g \ddot{\epsilon}_g^T) \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0} \\ &= \frac{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T \ddot{\epsilon}_g \ddot{\epsilon}_g^T \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0}{\sum_g \lambda_g^{CSS}} \\ &\quad + \frac{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T (H_g \ddot{\epsilon} \ddot{\epsilon}^T H_g^T - H_g \ddot{\epsilon} \ddot{\epsilon}_g^T - \ddot{\epsilon}_g \ddot{\epsilon}^T H_g^T) \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0}{\sum_g \lambda_g^{CSS}} \\ &= \Upsilon_1 + \Upsilon_2, \quad \text{where } \Upsilon_1 = \frac{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T \ddot{\epsilon}_g \ddot{\epsilon}_g^T \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0}{\sum_g \lambda_g^{CSS}}, \\ \Upsilon_2 &= \frac{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T (H_g \ddot{\epsilon} \ddot{\epsilon}^T H_g^T - H_g \ddot{\epsilon} \ddot{\epsilon}_g^T - \ddot{\epsilon}_g \ddot{\epsilon}^T H_g^T) \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0}{\sum_g \lambda_g^{CSS}} \end{aligned}$$

If the residuals $\hat{\epsilon}$ are close to the errors $\ddot{\epsilon}$, then Υ_2 will be small. One of the approximations applied by [Carter, Schnepel and Steigerwald \(2017\)](#) is to approximate Υ_0 as Υ_1 :

$$\begin{aligned} \Upsilon_0 &\approx \Upsilon_1 = \frac{c_0^T (\ddot{X}^T \ddot{X})^{-1} \left(\sum_g \ddot{X}_g^T \ddot{\epsilon}_g \ddot{\epsilon}_g^T \ddot{X}_g \right) (\ddot{X}^T \ddot{X})^{-1} c_0}{\sum_g \lambda_g^{CSS}} \\ \Upsilon_1 &= \frac{\sum_g \ddot{\epsilon}_g^T \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0 c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g \ddot{\epsilon}_g}{\sum_g \lambda_g^{CSS}} \end{aligned}$$

Let $\eta_g \sim N(0, I_g)$ such that $\ddot{\epsilon}_g = \sigma(1 - \rho)^{\frac{1}{2}} M_g \eta_g$:

$$\Upsilon_1 = \sigma(1 - \rho) \frac{\sum_g \eta_g M_g \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0 c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T M_g \eta_g}{\sum_g \lambda_g^{CSS}}$$

Notice that $\eta_g M_g \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0 c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T M_g \eta_g$ is a linear combination of χ_1^2 random variables with coefficients equal to the single non-zero eigenvalues of:

$$M_g \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0 c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T M_g$$

There is only one such non-zero eigenvalue, and that is:

$$c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T M_g \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0$$

Let $w_g \sim \chi_1^2$. It follows that:

$$\begin{aligned} \Upsilon_1 &\sim \sigma(1 - \rho) \frac{\sum_g w_g c_0^T (\ddot{X}^T \ddot{X})^{-1} \ddot{X}_g^T M_g \ddot{X}_g (\ddot{X}^T \ddot{X})^{-1} c_0}{\sum_g \lambda_g^{CSS}} \\ &\sim \frac{\sum_g w_g \lambda_g^{CSS}}{\sum_g \lambda_g^{CSS}} \end{aligned}$$

The second approximation applied (implicitly) by [Carter, Schnepel and Steigerwald \(2017\)](#) is the Satterthwaite approximation. Υ_1 is a linear combination of χ_1^2 random variables, and it is approximated by Υ^{CSS} , where $G^* \Upsilon^{CSS} \sim \chi_{G^*}^2$. Just like with Υ^{BM} above, the first two moments of Υ_1 and Υ^{CSS} match each other. So then let $\tau^{CSS} = \frac{z_0}{\sqrt{\Upsilon^{CSS}}}$ be the approximated test statistic with Υ^{CSS} substituted for Υ_0 . We have that $\tau^{CSS} \sim T(G^*)$.

Much of what I've ascribed to [Carter, Schnepel and Steigerwald \(2017\)](#) is implicit in the description of their test rather than explicit in their analysis. My purpose here was to explain what simplifications and approximations separate the reference distribution $T(G^*)$ from the exact distribution of the test statistic t_0 .